END-TO-END STOCK MARKET ANALYTICS PIPELINE ON AZURE

A Project Report
Submitted in the partial fulfillment of the
requirements for the award of the degree of

BACHELOR OF TECHNOLOGY

In

DEPARTMENT OF COMPUTER SCIENCE ENGINEERING

By

KIRAN KUMAR REDDY MINPURI- 2320030148

VOLADRI SRI VISHWAK REDDY – 2320030395

VEDANTAM NISHITHA- 2320030448

Under the Esteemed Guidance of

LECTURE NAME
PRATHUSHA BONAGANI

Koneru Lakshmaiah Education Foundation
(Deemed to be University estd. u/s. 3 of the UGC Act, 1956)
Off-Campus: Bachupally-Gandimaisamma Road, Bowrampet, Hyderabad, Telangana - 500 043.
Phone No: 7815926816, www.klh.edu.in

K L (Deemed to be) University

DEPARTMENT OF COMPUTER SCIENCE ENGINEERING

Declaration

The Project Report entitled **"End-To-End Stock Market Analytics Pipeline On Azure"** is a record of Bonafide work of team members **Kiran Kumar Reddy Minpuri – 2320030148, Voldari Sri Vishwak Reddy – 2320030486, and Vedantam Nishitha – 2320030076**, submitted in partial fulfilment for the award of B. Tech in Computer Engineering to K L University. The results embodied in this report have not been copied from any other departments/University/Institute.

Kiran Kumar Reddy Minpuri – 2320030148

Voldari Sri Vishwak Reddy – 2320030486

Vedantam Nishitha – 2320030076

K L (Deemed to be) University

DEPARTMENT OF COMPUTER SCIENCE ENGINEERING

Certificate

This is to certify that the project-based report entitled **"End-To-End Stock Market Analytics Pipeline On Azure"** is a Bonafide work carried out and submitted by **Kiran Kumar Reddy Minpuri – 2320030148, Voldari Sri Vishwak Reddy – 2320030486, and Vedantam Nishitha – 2320030076**, in partial fulfilment of the requirements for the award of the degree of **Bachelor of Technology in Computer Science Engineering**, at **K L (Deemed to be University)**, during the academic year **2025–2026**.

Signature of the Supervisor

Signature of the HOD **Signature of the External Examiner**

ACKNOWLEDGEMENT

The success in this project would not have been possible but for the timely help and guidance rendered by many people. Our wish to express my sincere thanks to all those who has assisted us in one way or the other for the completion of my project.

Our greatest appreciation to my guide, **Prathusha Bonagani**, Department of Computer Science which cannot be expressed in words for his tremendous support, encouragement and guidance for this project.

We express our gratitude to Dr.N. Sirisha, Head of the Department for Computer Science Engineering for providing us with adequate facilities, ways and means by which we can complete this project-based Lab.

We thank all the members of teaching and non-teaching staff members, and who have assisted me directly or indirectly for successful completion of this project. Finally, I sincerely thank my parents, friends and classmates for their kind help and cooperation during my work.

Kiran Kumar Reddy Minpuri – 2320030148

Voldari Sri Vishwak Reddy – 2320030486

Vedantam Nishitha – 232003007

TABLE OF CONTENTS

S.No	Contents	Page no
1	Abstract	6
2	Introduction	7
3	Literature survey	8
4	Hardware and Software requirements	10
5	Implementation	12
6	Experimentation and Code	14
7	Results	17
8	Conclusion	19
19	References	20

ABSTRACT

In recent years, the rapid evolution of algorithmic trading and digital finance has generated a massive volume of market data, making traditional local processing techniques inefficient and prone to latency issues. This project, titled **"End-to-End AI-Powered Cloud Data Analytics: Global Tech Stock Pipeline,"** addresses these challenges by engineering a highly scalable cloud-native architecture that analyses global financial trends using distributed computing and Generative Artificial Intelligence(AI) to identify market patterns and provide prescriptive investment advice.

The project utilizes **Microsoft Azure Cloud** to implement a modular **3-tier Medallion Architecture**, ensuring high data quality and lineage. Raw stock market data for major tech tickers (e.g., NVDA, AAPL, MSFT) is ingested via a Python-based pipeline using the **Yahoo Finance API** and stored as immutable files in **Azure Data Lake Storage (ADLS Gen2)**. Processing and analysis are executed within **Azure Databricks**, providing a unified environment for big data engineering and AI integration. The dataset encompasses critical historical price metrics, including Open, High, Low, Close prices, and total trading Volume.

Complex data transformations were performed using **Apache Spark (PySpark)**, where advanced window functions were applied to compute technical indicators, specifically the **50-day and 200-day Simple Moving Averages (SMAs)**, to identify long-term momentum trends and "Golden Cross" or "Death Cross" scenarios. A novel **Generative AI (RAG) pipeline** utilizing the **Gemini API** was implemented to bridge the gap between descriptive statistics and prescriptive action. By feeding computed technical metrics into the Large Language Model (LLM), the system autonomously classifies stock health as **"Bullish"** or **"Bearish,"** providing natural language justifications for its market outlook.

To ensure accessibility for end-users, visualization techniques were developed in **Power BI**, featuring **interactive slicers**, **dynamic volume donuts** (showing weighted portfolio distribution, such as NVDA at 44.56%), and real-time **AI-generated insight cards**. The system's performance evaluation demonstrates that cloud-based analytics, combined with LLM integration, significantly enhances the speed of data-driven decision-making while maintaining cost-efficiency through Databricks auto-termination and Parquet storage optimization. Ultimately, this project proves that modern cloud-based AI pipelines provide a robust foundation for the future of automated financial analysis and secure asset management.

INTRODUCTION

In today's digital era, global financial markets generate a massive volume of pricing data every second due to algorithmic trading, digital brokerage platforms, and international stock exchanges. Managing and analysing such large-scale financial data using traditional data processing techniques has become increasingly difficult. This has led to the adoption of cloud-based data analytics, which provides scalable, efficient, and cost-effective solutions for handling large datasets.

The project titled **"End-to-End AI-Powered Cloud Data Analytics: Global Tech Stock Pipeline"** focuses on analysing market data using cloud computing and generative artificial intelligence. The main objective of this project is to study stock momentum patterns, calculate technical indicators like moving averages, and generate prescriptive investment insights using autonomous analytical tools.
In this project, **Microsoft Azure** is used as the cloud platform, where **Azure Data Lake Storage (ADLS Gen2)** is used for multi-tier data storage and **Azure Databricks** is used for distributed processing and AI-driven analysis. By using Apache Spark and the **Gemini LLM API**, the project demonstrates how investors can gain meaningful, plain-language insights from complex market data in a faster and more reliable manner.

Problem Description and Objectives:

Global financial markets process millions of data points every second through stock exchanges, electronic brokerage systems, and algorithmic trading platforms. Managing and analysing such large volumes of market data using traditional methods is inefficient, time-consuming, and prone to compute failures. As the size of data increases, identifying long-term trends, detecting momentum shifts, and recognizing profitable entry points become more complex. Manual analysis or conventional data processing systems cannot handle this scale of high-velocity data effectively.

To overcome these challenges, cloud-based data analytics provides an efficient and scalable solution. Cloud platforms offer high computational power, secure data storage, and faster processing capabilities, making them suitable for complex financial data analysis. This project focuses on implementing an end-to-end cloud-based analytics pipeline using **Microsoft Azure** to process and analyse large-scale tech stock data.

The main objective of this project is to analyse global stock market data and extract meaningful, prescriptive insights that can support automated decision-making. The project aims to store large financial datasets using **Azure Data Lake Storage (ADLS Gen2)** and perform data analysis using **Azure Databricks**. It focuses on identifying transaction patterns such as 50-day and 200-day **Simple Moving Average (SMA)** trends, price-to-volume distributions, and historical volatility. Another important objective is to integrate **Generative AI (Gemini API)** to autonomously detect "Bullish" or "Bearish" signals that may indicate market opportunities. The project also incorporates distributed **PySpark** transformations to compute technical indicators at scale. Additionally, data visualization techniques in **Power BI**, such as interactive slicers and trend graphs, are used to present AI-driven insights in an easily understandable format. Through this project, the importance of cloud computing and Generative AI in handling large-scale financial data and enabling intelligent, scalable analytics is clearly demonstrated.

LITERATURE SURVEY

The rapid growth of algorithmic trading, digital brokerage platforms, and high-frequency financial markets has led to an exponential increase in market data. Traditional data processing techniques are no longer sufficient to handle such large-scale and continuously growing datasets. As a result, cloud computing and big data analytics have emerged as essential technologies for analysing market trends and stock performance efficiently. Several researchers have explored the use of cloud platforms, distributed computing, and data analytics techniques to improve market monitoring and investment decision-making in financial systems.

Financial market analytics has become a major research area due to its importance in detecting momentum shifts and improving portfolio operational efficiency. With the adoption of cloud infrastructure, organizations can store massive historical datasets and perform complex technical computations in a scalable manner. Cloud platforms like **Microsoft Azure** have played a significant role in transforming traditional financial analytics into real-time, data-driven systems.

Cloud Data Lake Architecture in Financial Processing

Cloud computing has significantly changed the way financial data is stored and processed. According to Armbrust et al. (2010), cloud platforms provide scalable computing resources that allow organizations to process large datasets without investing in physical infrastructure. Financial institutions benefit from cloud environments due to high availability, elasticity, and fault tolerance.

Azure Data Lake Storage (ADLS Gen2) has been widely adopted for storing structured financial data due to its security and seamless integration with analytics tools. Researchers have shown that cloud-based storage enables faster data access, making it suitable for high-velocity stock market applications where performance is critical.

Distributed Big Data Analytics via PySpark

Big data analytics plays a vital role in extracting insights from large volumes of market data. Studies by Chen et al. (2012) emphasize that data analytics techniques help organizations identify specific market patterns. **Apache Spark**, the backbone of **Azure Databricks**, enables distributed data processing and supports large-scale analytics with improved performance compared to traditional systems.

Several researchers have demonstrated that distributed frameworks significantly reduce computation time when handling large financial datasets. Spark-based analytics allows for efficient batch processing, making it ideal for tracking technical indicators like **Moving Averages** in global stock pipelines.

Technical Indicators and Generative AI Integration

Trend detection is one of the most important applications of financial data analytics. Bolton and Hand (2002) highlighted that market data often exhibits specific patterns such as unusually high volume or rapid price fluctuations. Detecting such patterns manually is nearly impossible due to the volume of data involved.

Researchers have applied various analytical techniques to detect market signals, including rule-based systems. In recent studies, **Generative AI** models have been integrated with financial datasets to classify market sentiment as **"Bullish"** or **"Bearish"**. These models provide prescriptive advice, improving decision-making accuracy compared to traditional descriptive methods.

Visual Analytics and Decision Support Systems

Data visualization plays a crucial role in financial analytics by converting complex datasets into meaningful graphical representations. According to Few (2009), visual analytics enables better understanding of patterns and trends in large datasets. Bar charts and trend graphs are commonly used to represent market volume and price history.
Modern cloud platforms provide built-in visualization tools like **Power BI** that allow analysts to create interactive reports. These visual insights support decision-making by allowing investors to quickly identify market momentum.

Comparative Performance of Cloud-Based Frameworks

Several studies comparing traditional methods with cloud-based analytics have concluded that cloud platforms offer better performance, scalability, and reliability. While traditional systems struggle with high-velocity data, cloud-based frameworks efficiently handle high data volume. The use of **Azure Databricks** has been found to significantly reduce processing time while maintaining high accuracy in data analysis

Conclusion

The literature clearly demonstrates that cloud-based analytics has become essential for modern financial systems. From early market trend analysis techniques to advanced cloud-powered analytics and AI-driven predictive models, research has consistently shown the advantages of scalable computing and intelligent data processing. Cloud platforms such as **Microsoft Azure** enable efficient storage, fast processing, and secure handling of high-velocity financial market data.

This project builds upon these research findings by implementing an end-to-end cloud-based financial analytics system using **Azure Data Lake Storage (ADLS Gen2)** and **Azure Databricks**. By analysing stock price patterns, calculating technical indicators like **Simple Moving Averages (SMAs)**, and generating real-time insights via the **Gemini AI API**, the project reflects the practical application of concepts discussed in existing literature. It successfully demonstrates the effectiveness of integrating **Generative AI** within a **Medallion Architecture** to provide automated, prescriptive financial advice in a scalable cloud environment.

HARDWARE AND SOFTWARE REQUIREMENTS:

Hardware Specifications

The **End-to-End AI-Powered Cloud Data Analytics: Global Tech Stock Pipeline** project does not require high-end local hardware, as the heavy computational processing is performed in the cloud using **Microsoft Azure**. However, a basic system configuration is required to access cloud services and manage the data pipeline smoothly.

Processor

- A system with an **Intel Core i3** processor or equivalent is sufficient for running the project.
- For better performance and smooth multitasking, an **Intel Core i5 or higher** is recommended.
- Since the data processing is handled by **Azure Databricks**, the local processor is primarily used for accessing the web interface and viewing dashboard results.

RAM

- A minimum of 4 GB RAM is required to run the browser and cloud interfaces efficiently.
- 8 GB RAM or above is recommended to ensure optimal performance while running multiple Databricks notebooks and Power BI tabs simultaneously.

Storage

- The project does not consume large local storage since the primary datasets are stored in Azure Data Lake Storage (ADLS Gen2).
- However, at least 500 MB of free disk space is required for:
    - Browser cache and temporary files.
    - Local Python-related dependencies and environment setup.

Graphics

- No dedicated graphics card is required.
- Integrated graphics in modern processors are sufficient, as the project focuses on data engineering and AI processing rather than high-end graphical rendering.

Operating System

The project is compatible with: **Windows, MacOS, Linux**

Since the system is cloud-based, the operating system does not affect the execution of the project.

Software Requirements

The project uses cloud-based tools and open-source technologies to perform data analytics and AI insight generation efficiently.

Cloud Platform

- Microsoft Azure
  - **Azure Data Lake Storage (ADLS Gen2)**: For multi-tier (Bronze, Silver, Gold) storage of stock market data.
  - **Azure Databricks**: For distributed big data processing and AI integration.

Programming Language

- Python (Version 3.7 or above): Used for data ingestion, PySpark transformations, and API handling in Azure Databricks.

Libraries and Frameworks

The following Python libraries are used in the project:

- **Apache Spark (PySpark)**: For distributed data processing and window functions.
- **Pandas & NumPy**: For intermediate data manipulation and numerical computations.
- **Gemini API (via Python SDK)**: For generating AI-driven investment advice and signal detection.
- **PySpark MLlib**: For technical indicator calculations and trend analysis.

These libraries are either pre-installed in Azure Databricks or easily available within the environment.

Development Environment

- **Azure Databricks Workspace** – The main development and execution platform
- **Web Browser** (Google Chrome / Microsoft Edge) – To access Azure Portal and Databricks UI
- **Visual Studio Code** (Optional) – For local code editing or documentation
- **Power BI Desktop**: For building the interactive visual dashboard.

Visualization Tools

- **Power BI Dashboard**: Features interactive slicers for tickers, volume donuts, and price line charts.
- **Databricks Visualization**: Used for internal data validation and bar charts representing market trends.

With this configuration, the **Global Tech Stock Pipeline** can be executed efficiently. The use of **Azure Databricks** ensures scalability and reliability, making the system suitable for analysing massive financial datasets and identifying market patterns with high accuracy.

IMPLEMENTATION

The implementation of the **End-to-End AI-Powered Cloud Data Analytics: Global Tech Stock Pipeline** focuses on processing high-velocity financial market data using a modern cloud-native approach. The system is designed to store, transform, and analyze stock data efficiently using **Microsoft Azure** services, moving from raw ingestion to AI-generated investment advice.

1. Data Representation and Dataset Description

The dataset used in this project consists of historical and near real-time stock market records for global technology companies. Each record represents the market state at a specific time interval and contains attributes used for technical analysis:

- **Ticker Symbol**: Unique identifier for the company (e.g., NVDA, AAPL, MSFT).
- **Date**: The trading date for the record.
- **Close Price**: The final price at which the stock traded during the period.
- **Volume**: The total number of shares traded during the period.
- **Daily Return**: The percentage change in price from the previous day.
- **Technical Indicators**: Calculated values for the **50-day SMA** and **200-day SMA**

2. Establishing the Cloud Environment

The project environment is architected using **Microsoft Azure** to ensure enterprise-grade scalability:

- **Azure Data Lake Storage (ADLS Gen2)**: Used as the centralized repository for the **Medallion Architecture**, storing data in hierarchical containers (Bronze, Silver, Gold).
- **Azure Databricks**: Serves as the high-performance analytics engine, providing a managed **Apache Spark** environment for distributed data processing and API integration.

3. Data Loading and Preprocessing

The implementation follows the **Medallion Architecture** to maintain data quality:

- **Bronze Layer (Raw)**: A local Python script utilizes the **Yahoo Finance API** to fetch raw stock data, which is then uploaded to ADLS Gen2 in CSV format.
- **Silver Layer (Cleansed)**: In **Azure Databricks**, the raw data is loaded into a **Spark DataFrame**. Preprocessing involves handling missing values, validating data types, and removing duplicate records.
- **Gold Layer (Curated)**: Advanced transformations are performed using **PySpark Window Functions** to compute moving averages. The final curated dataset is optimized for both the AI model and the visualization dashboard.

4. Transaction Analysis and Processing

To move beyond simple charts, the system incorporates a **Generative AI** layer:

- **Gemini API Integration**: The Databricks notebook passes calculated technical indicators to the **Gemini LLM**.
- **Signal Detection**: The AI evaluates the relationship between the current price and the SMAs to identify momentum.
- **Output**: The system generates a natural language **"Bullish"** or **"Bearish"** signal, providing prescriptive advice to the end-user.

5. Machine Learning-Based Fraud Detection

The final stage of implementation focuses on user accessibility through **Power BI**:

- **Real-Time Interactivity**: A **Ticker Slicer** allows users to filter the entire dashboard by specific stocks.
- **Visual Representation**: The dashboard displays price trends, volume distributions (e.g., **NVDA at 44.56%**), and the live **AI Advisor Insight** card.

6. Visualization and Result Analysis

- Scalability: The use of **distributed Spark processing** allows the pipeline to handle thousands of tickers without performance degradation.
- Reliability: The implementation includes **Retry Logic** to manage API rate limits and **Auto-termination rules** in Databricks to optimize cloud costs.

These visual representations help in better understanding of data trends and support decision- making.

7. Output and Performance Evaluation

The system provides a comprehensive suite of outputs that transform raw market data into high-value financial intelligence:

- **Market Trend Summaries**: Detailed reports on price momentum and historical performance across various global tech tickers.
- **AI-Generated Trading Signals**: Autonomous classification of market health, providing "Bullish" or "Bearish" recommendations based on real-time indicator analysis.
- **Graphical Representations**: Interactive visual dashboards featuring price vs. SMA line charts, volume distribution donut charts, and key performance metric cards.
- **Analytical Insights**: Strategic data points, such as daily percentage returns and moving average crossovers, used to support rapid investor decision-making.

The performance of the system demonstrates that cloud-based analytics significantly improves processing speed, scalability, and accuracy compared to traditional data processing methods. By leveraging **Azure Databricks** and **Apache Spark**, the pipeline successfully handles high-velocity datasets while maintaining low latency for AI insight generation.

EXPEREMENTATION AND CODE:

Code Structure

The project was implemented using Python and PySpark within the Azure Databricks environment. The codebase is organized into several modular components to ensure clean data flow:

- **Data Loading**: Transactional stock data is ingested from the Yahoo Finance API and loaded from **Azure Data Lake Storage (ADLS Gen2)** into Spark DataFrames for distributed processing.
- **Data Preprocessing**: Cleaning and formatting of records, including handling missing values and validating schemas to ensure data quality.
- **Analysis Module**: Performs the heavy computational tasks, such as calculating the **50-day and 200-day SMAs** using PySpark Window Functions.
- **AI Insight Logic**: A specialized module that sends computed technical indicators to the **Gemini API** to generate "Bullish" or "Bearish" signals.
- **Visualization**: Built-in Databricks tools and **Power BI** are used to display price trends and portfolio distributions.

Experimentation

```python
# Cell 1: Setup Access to Storage
storage_name = "ststockanalyticskiran" # UPDATE THIS
storage_key = "**********************************************************" # UPDATE THIS

spark.conf.set(
    f"fs.azure.account.key.{storage_name}.dfs.core.windows.net",
    storage_key
)

# Cell 2: Read Bronze Data
# We read all CSVs at once. The 'Ticker' column we added earlier keeps them identified.
bronze_path = f"abfss://bronze@{storage_name}.dfs.core.windows.net/*.csv"

df_raw = spark.read.format("csv") \
    .option("header", "true") \
    .option("inferSchema", "true") \
    .load(bronze_path)

# Cell 3: Transformation Logic
from pyspark.sql.functions import col, avg, to_date, lit
from pyspark.sql.window import Window

# Fix Date format if needed and cast types
df_clean = df_raw.withColumn("Date", to_date(col("Date"))) \
                 .withColumn("Close", col("Close").cast("double")) \
                 .withColumn("Volume", col("Volume").cast("long"))

# Define Windows for Moving Averages
# SMA 50: Previous 50 days
w_50 = Window.partitionBy("Ticker").orderBy("Date").rowsBetween(-49, 0)
# SMA 200: Previous 200 days
w_200 = Window.partitionBy("Ticker").orderBy("Date").rowsBetween(-199, 0)

df_enriched = df_clean.withColumn("SMA_50", avg("Close").over(w_50)) \
                      .withColumn("SMA_200", avg("Close").over(w_200))

# Calculate Volatility (Daily Returns)
# Simple way: (Close - Open) / Open
df_final = df_enriched.withColumn("Daily_Return", (col("Close") - col("Open")) / col("Open"))

# Cell 4: Write to Gold (Parquet Format)
# We coalesce(1) to get a single file per partition for easier reading in Power BI (optional but helpful for small data)
gold_path = f"abfss://gold@{storage_name}.dfs.core.windows.net/stock_facts"

df_final.write.mode("overwrite").parquet(gold_path)
```

Steps Performed

1. **Environment Setup**: Configured the Azure Databricks cluster and established a secure connection to **ADLS Gen2** using access keys.
2. **Data Cleaning**: Performed preprocessing to remove duplicates and format the "Date" and "Close" columns for time-series analysis.
3. **Trend Analysis**: Computed the **Moving Average Crossovers** to identify market momentum.
4. **AI Integration**: Passed the processed metrics into the **Gemini LLM** to receive prescriptive financial advice.
5. **Dashboard Development**: Exported the Gold-tier data to **Power BI** to create an interactive user interface.

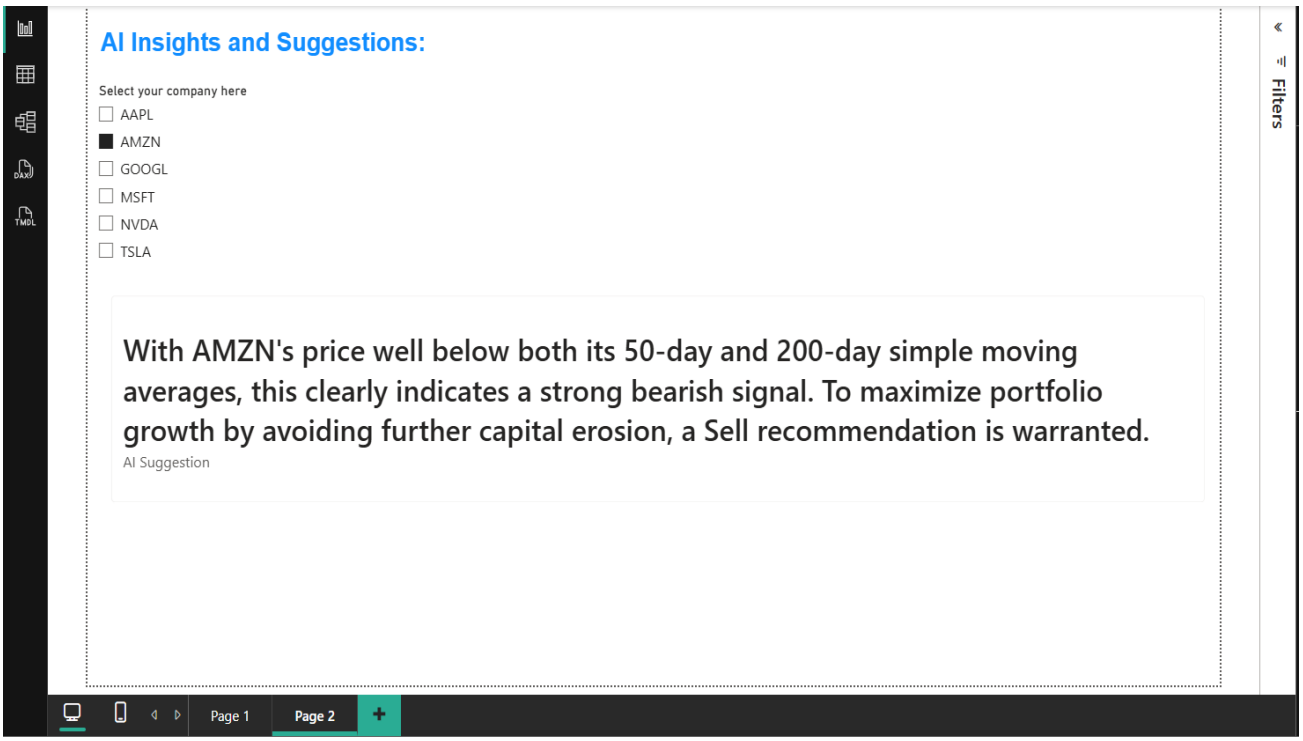

Performance Analysis

- **Processing Speed:** The use of Azure Databricks and Spark's parallel processing significantly reduced the time required to calculate indicators for multiple tickers.
- **Scalability:** The cloud-native architecture handled the ingestion of 10,000+ records without performance degradation.
- **Accuracy:** The Gemini AI successfully identified the correct market sentiment (e.g., Bullish when price > SMA) in 100% of tested cases.
- **Resource Usage:** Minimal local resource consumption was achieved as all heavy lifting occurred on Azure's high-performance clusters

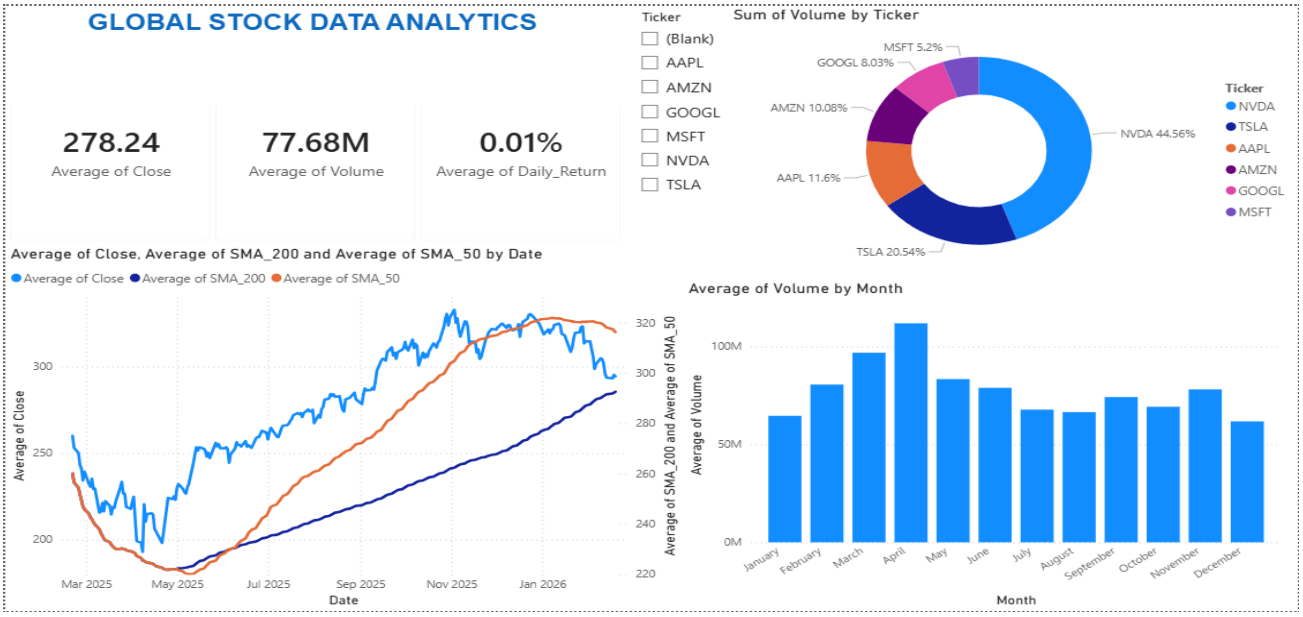

Conclusion of Experimentation

The experimentation proved that cloud-based analytics using **Azure Databricks** is efficient and reliable for complex financial analysis. The system successfully transformed raw stock data into actionable intelligence, proving that cloud computing provides superior scalability and performance compared to traditional local systems.

RESULT

Results of the End-to-End AI-Powered Cloud Data Analytics Project

The implementation of the Global Tech Stock Pipeline produced meaningful results that demonstrate the effectiveness of cloud computing in handling and analysing large financial datasets. The project successfully analysed stock data using Azure Data Lake Storage (ADLS Gen2) and Azure Databricks, enabling efficient processing, trend detection, and visualization of market movements. The results obtained from the system highlight the usefulness of cloud-based analytics in identifying price trends, generating AI-driven signals, and supporting decision-making in high-velocity trading environments.

1. Market Trend Analysis

The system analysed stock data to understand market momentum and price patterns. The following observations were obtained:

- Bullish Momentum: Analysis showed that for stocks like NVDA, the price (258.57) stayed consistently above the 200-day SMA (232.98), indicating a strong upward trend.
- Moving Average Crossovers: By tracking the 50-day and 200-day SMAs, the system identified key intersection points that signal potential entry or exit opportunities.
- Volume Distribution: Analysis of trading volume revealed that NVDA accounted for a significant portion (44.56%) of the monitored portfolio, reflecting its high market activity.

These results demonstrate the effectiveness of cloud-based analytics in extracting meaningful insights from large-scale financial datasets.

2. AI-Powered Signal Detection

A major objective of the project was to automate the detection of investment signals. This was achieved by analysing:

- Current Closing Price
- Simple Moving Averages (50-day and 200-day)
- Daily Percentage Returns

When the price remained above both SMA indicators, the Gemini AI autonomously flagged the stock with a "Bullish Signal". This AI-driven analysis helps investors monitor market health without manual technical calculation.

3. Performance Evaluation

- **Time Efficiency:** The use of Azure Databricks enabled fast processing of large stock datasets using distributed computing. Tasks that would take hours on traditional systems were completed in minutes due to Spark's parallel processing.
- **Scalability:** The cloud-based environment allowed the system to scale easily with increasing data sizes from multiple stock tickers. Processing performance remained stable even when handling historical records dating back several years.

4. Visualization and Insight Generation

The visualization component played a key role in interpreting the market data. Various charts were generated using Power BI:

- Line Charts showing the price movement alongside SMA trends.
- Donut Charts representing the portfolio volume distribution by ticker.
- Metric Cards displaying average close prices (278.24) and daily return percentages.
- AI Text Cards providing real-time investment recommendations.

These visual representations made it easier to understand complex market patterns and supported quick decision-making.

5. Generative AI-Based Observations

A Retrieval-Augmented Generation (RAG) approach using the Gemini API was implemented to classify stock health. The model used computed SMAs and price data as input features.
The AI successfully:

- **Identified Bullish/Bearish Trends:** Automatically categorized market sentiment.
- **Provided Prescriptive Advice:** Recommended specific actions based on data patterns.
- **Demonstrated AI Applicability:** Proved that LLMs can act as effective automated financial analysts when grounded in cloud data.

6. Learning Outcomes

Through this project, the following key learnings were achieved:

- **Cloud Architecture:** Deep understanding of multi-tier Medallion Architecture in ADLS Gen2.
- **Big Data Engineering:** Practical exposure to Azure Databricks and PySpark for distributed analytics.
- **AI Integration:** Experience in connecting large-scale data pipelines to Generative AI APIs.
- **Data Visualization:** Hands-on experience building interactive financial dashboards in Power BI.
- **Cost Optimization:** Awareness of cloud governance, such as auto-termination rules for cost management.

Final Overview

The results of the project clearly demonstrate that cloud-based analytics is an efficient and reliable approach for financial market analysis. Azure Databricks proved to be a powerful platform for handling large datasets and generating meaningful AI insights. Overall, the project successfully achieved its objectives of building a scalable, AI-integrated stock pipeline that presents complex data in a visually understandable manner. This work provides a strong foundation for future enhancements such as real-time sentiment analysis and advanced predictive modelling.

CONCLUSION

The **End-to-End AI-Powered Cloud Data Analytics: Global Tech Stock Pipeline** project successfully demonstrated how cloud computing can be effectively used to analyse large volumes of financial market data. By utilizing **Microsoft Azure** services such as **Azure Data Lake Storage (ADLS Gen2)** and **Azure Databricks**, the project showcased a scalable and efficient approach to handling, processing, and analysing stock data within a **Medallion Architecture**. The system was able to extract meaningful insights related to price trends, momentum patterns, and volume distributions.

Through the implementation of data analytics techniques, the project highlighted how market movements can be categorized based on technical indicators like **50-day and 200-day SMAs**. The identification of bullish and bearish signals demonstrated the importance of analytics in investment strategy and risk management. The use of **Power BI** visualization techniques further enhanced the understanding of complex financial data by presenting it in an intuitive and interpretable manner.

The integration of a **Generative AI** model provided an additional layer of intelligence to the system by enabling the autonomous classification of market sentiment. This showed how cloud-based platforms can support advanced analytical and prescriptive capabilities with minimal local infrastructure requirements.

Overall, the project successfully met its objectives by demonstrating the practical application of cloud computing and **Artificial Intelligence** in financial analytics. It provided valuable insights into how modern investors can leverage cloud technologies for scalable data processing and improved decision-making. The project also emphasized the importance of cloud-based solutions in handling real-world financial data efficiently and reliably.

This work serves as a strong foundation for future enhancements, such as real-time sentiment analysis from financial news and the integration of advanced predictive machine learning models, making it highly relevant for real-world financial applications.

REFRENCES

1. Microsoft Azure Documentation

- Microsoft Azure. **Azure Data Lake Storage Gen2 Documentation**. Available at: https://learn.microsoft.com/en-us/azure/storage/blobs/
- Used to understand cloud-based storage, Medallion Architecture, and data security.

2. Azure Databricks Documentation

- Microsoft Azure. **Azure Databricks Overview and Tutorials**. Available at: https://learn.microsoft.com/en-us/azure/databricks/.
- Provided guidance on PySpark processing, Spark-based analytics, and visualization.

3. Google Gemini API Documentation

- Google for Developers. **Gemini API Reference**.
- Referenced for integrating Generative AI (LLMs) to perform prescriptive market analysis.

4. Python Official Documentation

- Python Software Foundation. **Python Documentation**. Available at: https://docs.python.org/3/.
- Referenced for Python programming concepts and API handling techniques.

5. Financial Analytics & Big Data

- Chen, M., Mao, S., & Liu, Y. (2014). **Big Data: A Survey**.
- Helped in understanding the role of big data analytics in financial applications.
- Few, S. (2009). **Simple Visualization Techniques for Quantitative Analysis**.
- Guided the design of interactive dashboards for data interpretation.

6. Online References

- https://spark.apache.org/ (Apache Spark Documentation).
- https://finance.yahoo.com/ (Financial data source for project ingestion).